




PHASE 1: Semantics-Aware Head Identification



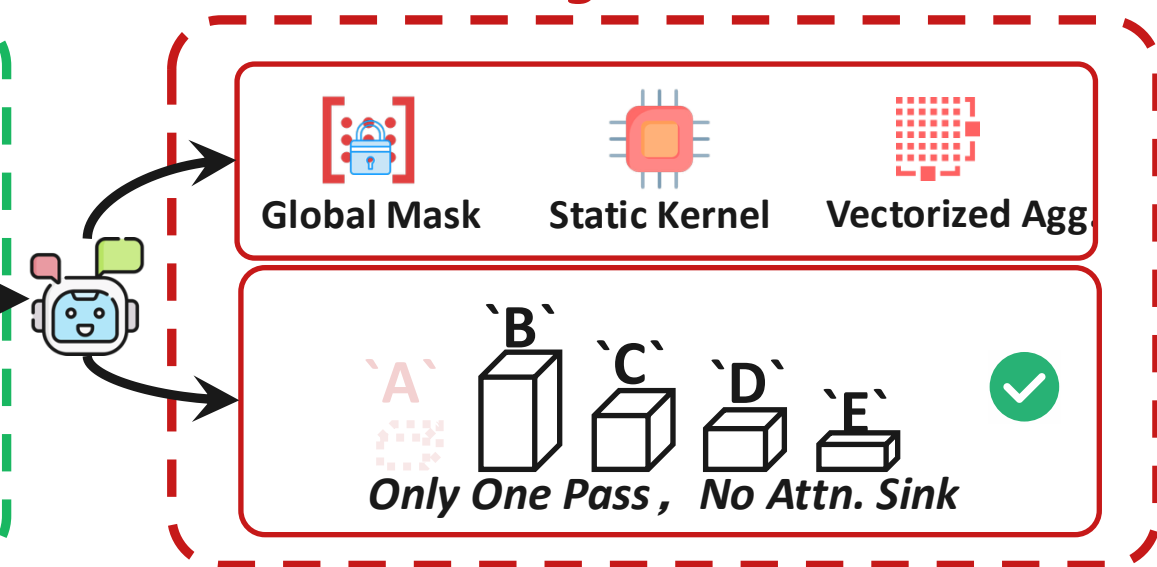
PHASE 2: Dynamic Pruning



PHASE 3: Calibration-Free SFT via Bias Suppression



PHASE 4: Single-Pass Inference



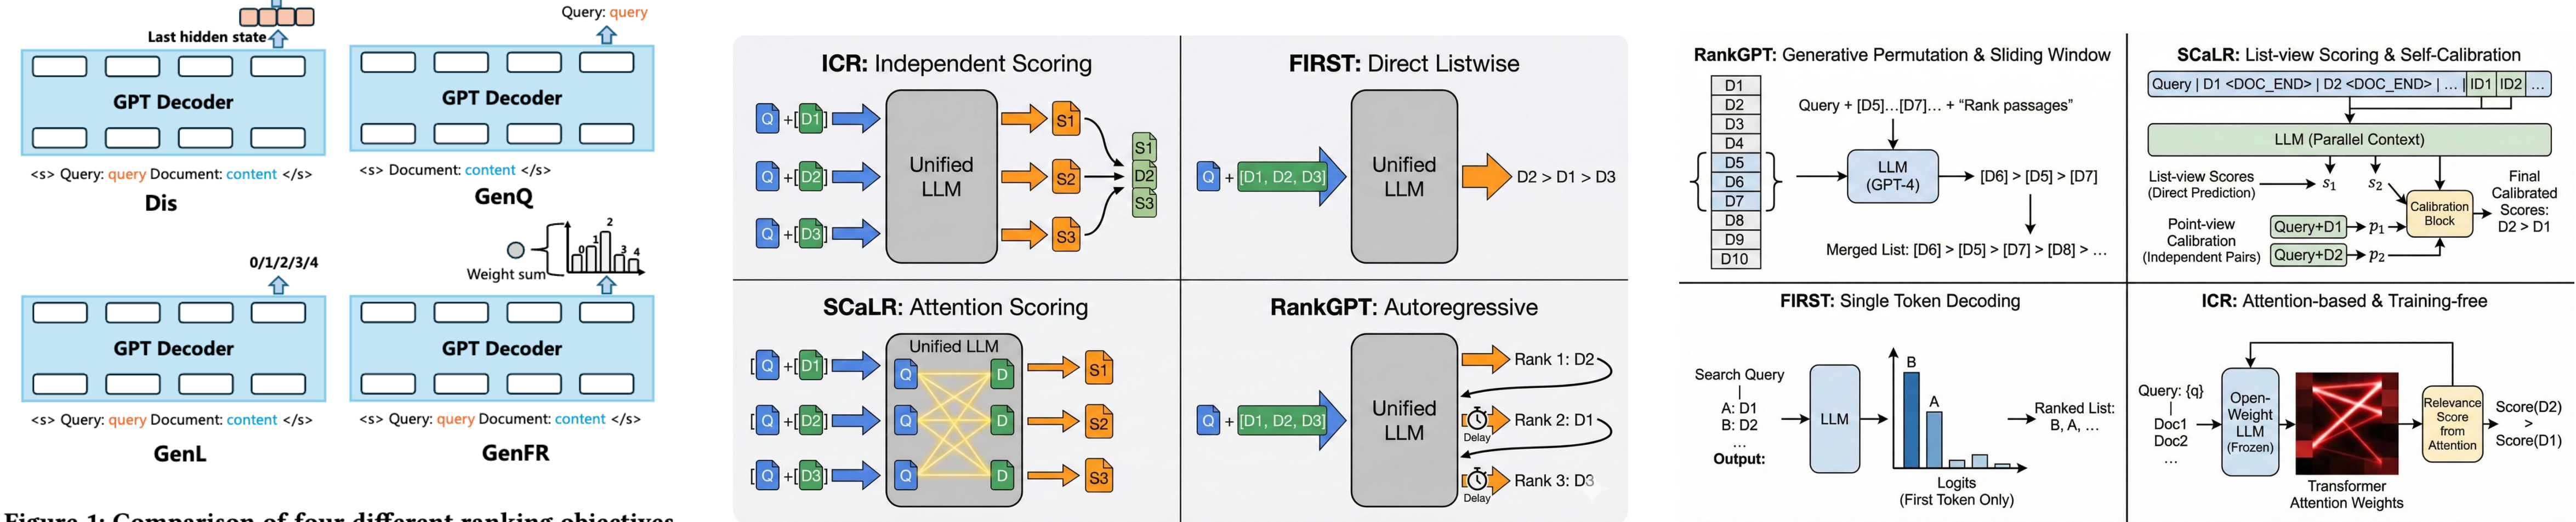


Figure 1: Comparison of four different ranking objectives for large language models.

